Big Data and Security

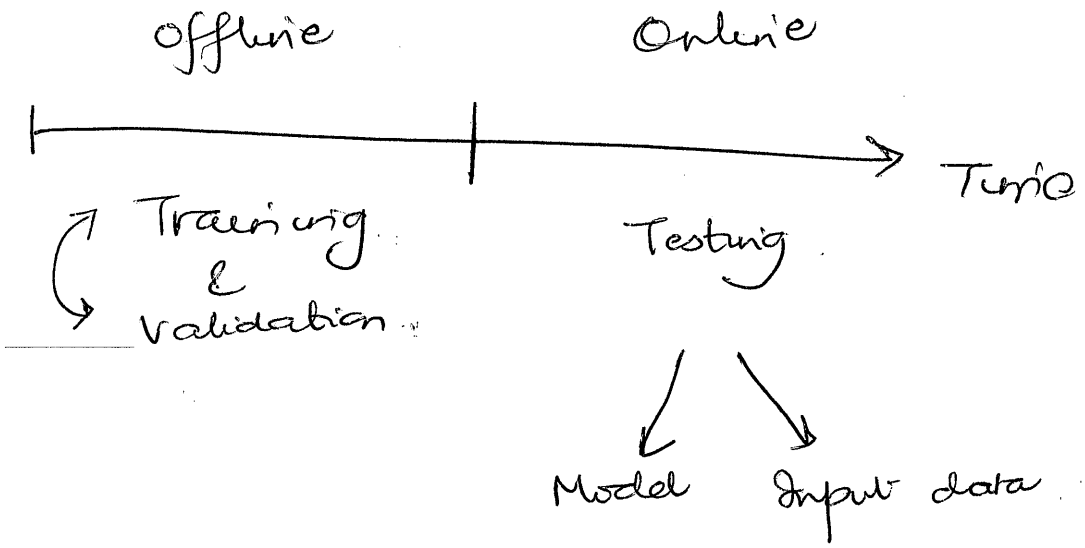

Black-box attack

White-box attack

$f(x) \rightarrow$ Malicious

Adversary creates

$x' = x + \Delta$ s.t. $f(x') \rightarrow$ Benign.

$f(x)$: [0.2 (Benign) 0.8 (Malicious)] $f(x')$: [0.4 (Benign) 0.6 (Malicious)]

Logit vector

$x \longrightarrow x' = x + \epsilon$

$f(x)$ Real value $f(x')$

Gradient $\lim_{\epsilon \to 0} \frac{|f(x') - f(x)|}{\epsilon}$

Defenses against Evasion Attacks

Adversarial Training

Original training data

D

$D_i \xrightarrow{\text{Perturbation}} D_i'$

Model being trained to approximate fn f

$$f(D_i) = f(D_i')$$

Perturbation function: Generative Adversarial Network (GAN)

Defensive Distillation

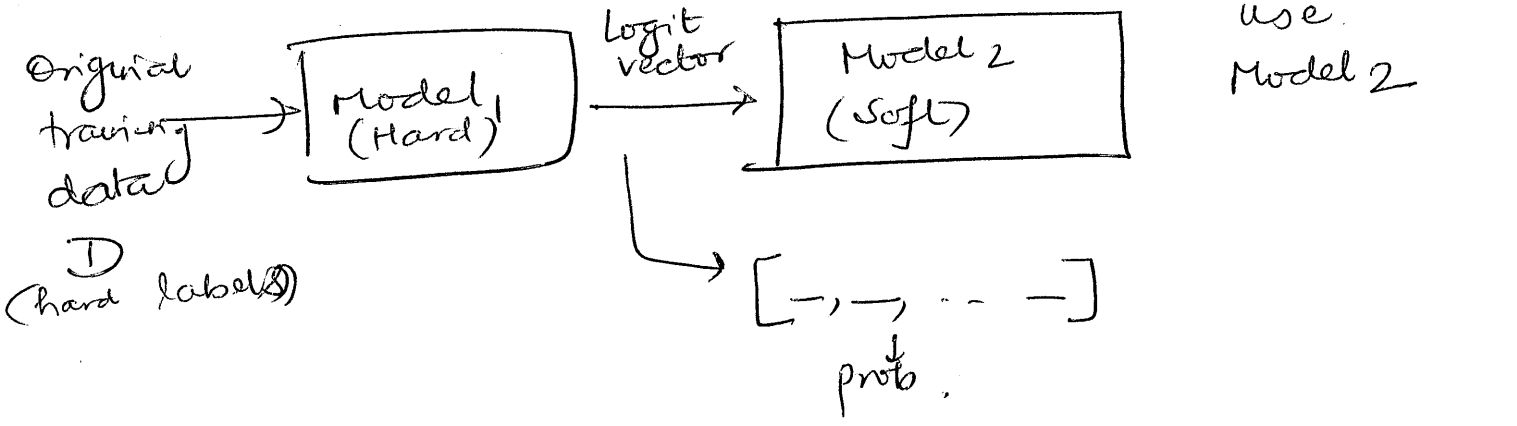

$\text{Prediction}_0 = [a_1, a_2, a_3, a_4, a_5]$

$\text{Prediction}_1 = [b_1, b_2, b_3, b_4, b_5]$

$L_{\#1}\ \text{norm} = \sum_{i=1}^{5} |a_i - b_i|$

Model : f

$f(D_i)$

$f(D_i\text{-squeezed})$

$d_1 \leftarrow L_1\ (f(D_i), f(D_i\text{-squeezed}))$

if $(d_1 > T)$ then D_i is adversarial

else D_i is legitimate

$T \uparrow$ False positive $\downarrow$ Good

Missed detection $\uparrow$ Bad