

ECE 60872/CS 590: Fault-Tolerant Computer System Design Software Fault Tolerance

Saurabh Bagchi
School of Electrical & Computer Engineering
Purdue University



Some material based on ECE442 at the University of Illinois taught by
Prof. Ravi Iyer & Zbigniew Kalbarczyk

Outline

- **Definition and Motivation for Software Fault Tolerance**
- **Process pairs**
- **Robust data structures**

What is Software Fault Tolerance?

- Three alternative definitions
 1. Management of faults originating from defects in design or implementation of software components
 2. Management of hardware failures in software
 3. Management of network failures
- We will follow the classical definition (1) due to Avizienis in 1977

Motivation for Software Fault Tolerance

- Usual method of software reliability is fault avoidance using good software engineering methodologies
- Large and complex systems \Rightarrow fault avoidance not successful
 - Rule of thumb fault density in software is 10-50 per 1,000 lines of code for good software and 1-5 after intensive testing using automated tools
- Redundancy in software needed to detect, isolate, and recover from software failures
- Hardware fault tolerance easier to assess
- Software is difficult to prove correct

HARDWARE FAULTS

1. Faults time-dependent
2. Duplicate hardware detects
3. Random failure is main cause

SOFTWARE FAULTS

- Faults time-invariant
Duplicate software not effective
Complexity is main cause

Consequences of Software Failure

- General Accounting Office reports \$4.2 million lost annually due to software errors
- Launch failure of Mariner I (1962)
- Destruction of French satellite (1988)
- Problems with Space Shuttle and Apollo missions
- SS7 (signaling system) protocol implementation - untested patch (mistyped character) (1997)
- Therac 25 (overdose of medical radiation 1000's of rads in excess of prescribed dosage)
- Toyota Prius recall (2004) due to bug in embedded code

Purdue University
60872/CS 590

5

ECE

Difficulties

- Improvements in software development methodologies reduce the incidence of faults, yielding fault avoidance
- Need for test and verification
- Formal verification techniques, such as proof of correctness, can be applied to rather small programs
- Potential exists of faulty translation of user requirements
- Conventional testing is hit-or-miss. "Program testing can show the presence of bugs but never show their absence," - Dijkstra, 1972.
- There is a lack of good fault models for software defects

Purdue University
60872/CS 590

6

ECE

Forms of Software Testing

- Exhaustive testing of reasonable sized applications is impossible
- Approach is to define equivalence classes of inputs so that only one test case from each class suffices
- Techniques proposed include
 - Path testing
 - Branch testing
 - Interface testing
 - Special values testing
 - Functional testing
 - Anomaly analysis
- Studies have shown path testing and interface testing while difficult to design afford good coverage for large number of applications

Approaches to Software Fault Tolerance

- **ROBUSTNESS:** The extent to which software continues to operate despite introduction of invalid inputs.
Example: 1. Check input data
 - =>ask for new input
 - =>use default value and raise flag2. Self checking software
- **FAULT CONTAINMENT:** Faults in one module should not affect other modules.
Example: Reasonable checks
Watchdog timers
Overflow/divide-by-zero detection
Assertion checking
- **FAULT TOLERANCE:** Provides uninterrupted operation in presence of program fault through multiple implementations of a given function

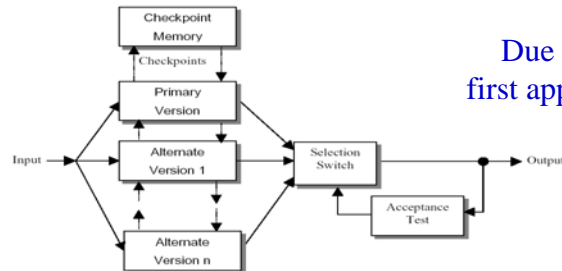
Temporal Redundancy

- Reexecution of a program when error is encountered
- Error may be faulty data, faulty execution or incorrect output
- Reexecution will clear errors arising from temporary circumstances
- Examples: Noisy communication channel, Full buffers, Power supply transients, Resource exhaustion in multiprocess environment
- Provides fault containment
- Possible to apply to applications with loose time constraints

Multi-Version Software Fault Tolerance

- Use of multiple versions (or “variants”) of a piece of software
- Different versions may execute in parallel or in sequence
- Rationale is that multiple versions will fail differently, i.e., for different inputs
- Versions are developed from common specifications
- Three main approaches
 - Recovery Blocks
 - N-version Programming
 - N Self-Checking Programming


Recovery Blocks



Due to Brian Randell,
first appeared in ToSE 1975

- Checkpoint and restart approach
 - Try a version, if error detected through acceptance test, try a different version
 - Ordering of the different versions according to reliability
- Checkpoints needed to provide valid operational state for subsequent versions
- Acceptance test could be on output or embedded in code

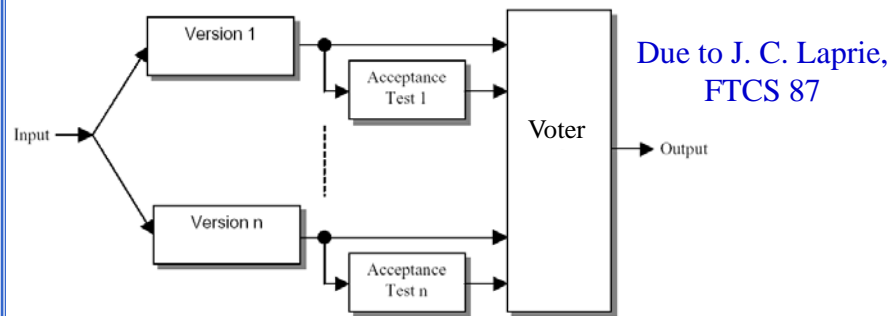
N-Version Programming



Due to Al Avizienis,
first appeared in CompSAC 1977

- All versions designed to satisfy same basic requirement
- Decision of output comparison based on voting
- Different teams build different versions to avoid correlated failures

N Self-Checking Programming





- Multiple software versions with structural variations of RB and NVP
- Use of separate acceptance tests for each version



Reliability Analysis of Multi-Version Approaches

- Three postulates of software development:
 - P1: Complexity Breeds Bugs: Everything else being equal, the more complex the software project is, the harder it is to make it reliable.
 - P2: All Bugs are Not Equal: You fix a bunch of obvious bugs quickly, but finding and fixing the last few bugs is much harder, if you can ever hunt them down.
 - P3: All Budgets are Finite: There is only a finite amount of effort (budget) that we can spend on any project. That is, if we go for n version diversity, we must divide the available effort n-way.
- $R(t) = e^{-\lambda t}$
- Failure rate $\lambda \propto 1/\text{Effort (E)}$
- Failure rate $\lambda \propto \text{Complexity (C)}$

Reliability of NVP vs. single version



Reliability of RB vs. Simplex




Process Pairs

- Used in HP Himalaya servers as part of their NonStop Advanced Architecture
- Bragging rights of the architecture
 - Run the majority of credit and debit card systems in N.America
 - More than US\$3 billion of electronic funds transfers daily
 - Run many of the E911 systems in North America
- Primary and backup processes on two different processors
- Primary process executes actively
 - Backup process is kept current by periodically sending state of primary process
- Processors execute fail-stop failure
 - When processor failure detected, backup takes over

Process Pairs

- **Applicability**
 - Permanent and transient hardware and software failures
 - Loosely coupled redundant architectures
 - Message passing process communication
 - Well suited for maintaining data integrity in a transactional type of system
 - Can be used to replicate a critical system function or user application
- **Assumptions**
 - Hardware and software modules design to *fail-fast*, i.e., to rapidly detect errors and subsequently terminate processing
 - Errors can be corrected by re-executing the same software copy in changed environment

Process Pairs Mechanism in Tandem Guardian OS



1. The application executes as *Primary*
2. *Primary* starts a *Backup* on another processor
3. Duplicated file images are also created
4. *Primary* periodically sends checkpoint information to *Backup*
5. *Backup* reads checkpoint messages and updates its data, file status, and program counter
 - the checkpoint information is inserted in the corresponding memory locations of the *Backup*
7. *Backup* loads and executes if the system reports that *Primary* processor is down
 - the error detection is done by *Primary* OS or
 - *Primary* fails to respond to "I am alive" message
8. All file activities by *Primary* are performed on both the primary and backup file copies
9. *Primary* periodically asks the OS if a *Backup* exists
 - if there is no *Backup*, the *Primary* can request the creation of a copy of both the process and file structure

Evaluation of Process-Pairs

- Done for Tandem's Guardian OS Studied Tandem Product Report (TPR) which are used to report product failures
- Problem classified as software fault only after analysts have pinpointed the cause
- Classes of software faults (not exhaustive)
 - Incorrect computation (3%)
 - Data fault (15%)
 - Missing operation (20%)
 - Side effect of code update (4%)
 - Unexpected situation (29%)
 - Microcode defect (4%)

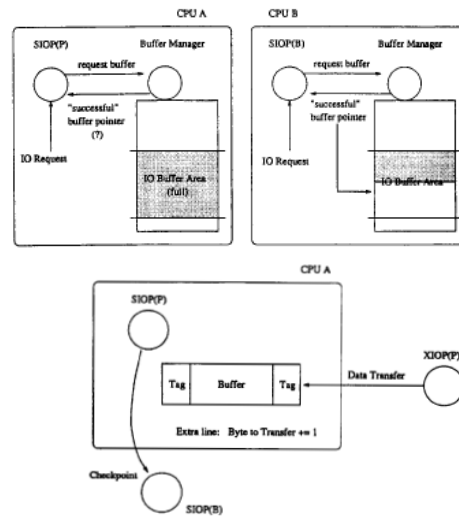
Results from Evaluation

- Out of total software failures, 138 out of 169 (82%) caused single processor halt (recoverable). This is a measure of the software fault tolerance of the system.
- Reasons for multiple processor fault
 - Same fault as in the primary: 17/28 (60%)
 - Second fault during job recovery: 4/28 (14.3%)
 - Second halt is not related to process pairs: 4/28 (14.3%)

Results from Evaluation

- Reasons for uncorrelated software fault
 - Backup reexecutes same task, but same fault not exercised: 29%.
 - Different memory state
 - Race or timing related problem
 - Example:
 - Privileged process on primary requests a buffer
 - Because of high user activity on primary, buffer exhaustion
 - Bug in buffer management routine and returns “success”
 - Primary privileged process uses uninitialized buffer pointer and causes processor halt
 - Backup process served the request after takeover
 - But buffer was available on the backup processor

Figure for Cases of Software Fault Tolerance



Results from Evaluation

- **Reasons for uncorrelated software fault**
 - Backup does not reexecute failed request on takeover: 20%.
 - Processor monitoring task
 - Interactive task
 - Effect of error latency: 5%
 - Task that caused the error finished before detection
 - Example: I/O process for copying buffer from source to destination.
 - Copied an additional byte overwriting buffer tag.
 - No problem in data transfer.
 - The successful data transfer was checkpointed but not the corrupted buffer tag
 - Problem surfaces later when buffer manager verifies buffer.
 - No problem when reexecuting on backup.

Results from Evaluation

- Process pairs with checkpointing and restart recovers from 75% of reported software faults that result in processor failures
- The complexity of process pairs introduces some faults
 - 16% of single processor halts were failures of backup processes
- Counter-intuitive result since same software run on both processors
- Loose coupling between processors, long error latency, operation using checkpoints and not lock-step
- Are process triples better than process pairs?

Process Pairs Advantages & Disadvantages

Advantages

- Extremely successful in Tandem OLTP applications
- Tolerates hardware, operating system, and application failures
- High coverage (> 90%) of hardware and software faults
- The backup does not significantly reduce the performance

■ Disadvantages

- Necessity of error detection checks and signaling techniques to make a process fail-fast
- Process pairs are difficult to construct for non-transaction-based applications

Robust Data Structures

- The goal is to find storage structures that are robust in the face of errors and failures
- What do we want to preserve?

Semantic integrity - the data meaning is not corrupted

Structural integrity - the correct data representation is preserved

Focus on techniques for preserving the structural integrity

Robust Data Structures (cont.)


- A robust data structure contains *redundant data* which allow *erroneous changes* to be detected, and possibly corrected
 - a change is defined as an elementary (e.g., as single word) modification to the encoded (data structure representation on a storage medium) form of a data structure instance
 - structural redundancy
 - a stored count of the numbers of nodes in a structure instance
 - identifier fields
 - additional pointers

Robust Data Structures (cont.)

- Consider data structure which consists of **a header** and a set of **nodes**
 - the header contains
 - pointers to certain nodes of the instance or to parts of itself
 - counts
 - identifier fields
 - a node contains
 - data items
 - structural information: pointers and node type identifier fields
- Error detection and correction
 - in-line checks may be introduced into normal system code to perform error detection and possibly correction, during regular operation

Link Lists

- Non-robust data structure
 - in each node store a pointer to the next node of the list
 - place a null pointer in the last node




0-detectable and 0-correctable
changing one pointer to NULL can
reduce any list to empty list

Robust Data Structures

Single-Linked List Implementation

– Additions for improving robustness

- an identifier field to each node
- replace the NULL pointer in the last node by a pointer to the header of the list
- stores a count of the number of nodes



1-detectable and 0-correctable


- change to the count can be detected by comparing it against the number of nodes found by following pointers
- change to the pointer may be detected by a mismatch in count number or the new pointer points to a foreign node (which cannot have a valid identifier)

Robust Data Structures

Double-Linked List Implementation

■ Additions for improving robustness

- a pointer added to each node, pointing to the predecessor of the node on the list

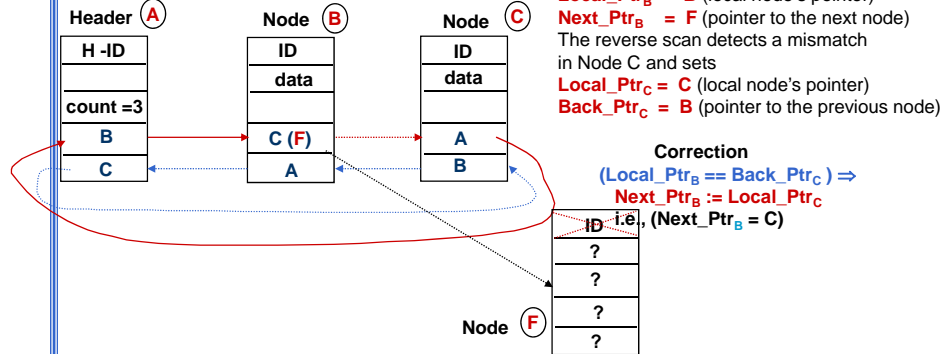


2-detectable and 1-correctable

the data structure has two independent, disjoint sets of pointers, each of which may be used to reconstruct the entire list

Error Correcting in Double-Linked List

- Scan the list in the forward direction until an identifier field error or forward/backward pointer mismatch is detected
- When this happens scan the list in the reverse direction until a similar error is detected
- Repair the data structure



Application of Robust Data Structures for Semantic Error Checking

- Application to checking index corruption in B-trees
- See class presentation

Robust Data Structures

Concluding Remarks

- Commonly used techniques for supporting robust data structures
 - techniques which preserve structural integrity of data
 - binary trees, heaps, fifos, queues, stacks
 - linked data structures
 - content-based techniques
 - checksums, encoding
- Limitations
 - not transparent to the application
 - best in tolerating errors which corrupt the structure of the data (not the semantic)
 - increased complexity of the update routines may make them error prone
 - erroneous changes to the data structure may be propagated by correct update routines
 - faulty update routines may provoke correlated erroneous changes to several fields

References

- "History's Worst Software Bugs" Wired Magazine, 11/8/05.
- D. K. Pradhan, ed., "Fault Tolerant Computer System Design", Chapter 7: Fault Tolerance in Software
- Multi-version software
 - Lui Sha, "Using Simplicity to Control Complexity," IEEE Software, Jul/Aug 01, pp. 20-28.
 - Wilfredo Torres-Pomales, "Software Fault Tolerance: A Tutorial," Technical Report: NASA-2000-tm210616, 2000.
- Process pairs
 - Inhwon Lee, R.K. Iyer: "Software dependability in the Tandem GUARDIAN system", IEEE Transactions on Software Engineering, May 1995.
- Robust data structures
 - David J. Taylor, David E. Morgan, James P. Black: "Redundancy in Data Structures: Improving Software Fault Tolerance." TSE 6(6): 585-594 (1980)
 - K. Fujimura, P. Jalote, "Robust search methods for B-trees", FTCS-18, June 1988.